

A Generalized Model of Consistency and Isolation Levels

Subhajit Sidhanta
INESC-ID/IST

Ricardo Dias
SUSE Linux GmbH

Rodrigo Rodrigues
INESC-ID/IST

ABSTRACT

The main contribution of this paper is a specification language *ConSpec*, that enables formalization of any possible consistency model or isolation level that a storage system may provide, in a uniform syntax and form, irrespective of the design and implementation of the target storage system. Prior to ConSpec, definitions of consistency models were tightly coupled with the specific design and implementation of the target storage system. Each new storage system would redefine a consistency model in terms of its design and implementation. The absence of a uniform definition makes it impossible to compare and verify the consistency provided by a given target storage system. Existing definitions of consistency models and isolation levels either use textual descriptions or graphical representation of proscribed anomalies and allowed dependencies between operations, making automated verification even more difficult. To the best of our knowledge, this is the first work that expresses both consistency and isolation levels in terms of ConSpec, a language that extends Linear Temporal Logic (LTL). Use of LTL ushers in the possibility of leveraging existing tools like satisfiability solvers and theorem provers to build automated verifiers. Finally, using insights derived from our analysis, we restate the CAP theorem to take into account the plethora of possible consistency models, in contrast with the original definition which only considered linearizability.

1. INTRODUCTION

Development of modern internet-based applications and services requires application developers to take into account the consistency [4] and isolation level [1] provided by the underlying storage system, to ensure correctness and performance guarantees. Currently, there is a large number of possible consistency models and isolation levels, many of which are not clearly defined using formal semantics. The universe of consistency models and isolation levels is, in fact, expanding constantly. Newer consistency models are coming up regularly to address specific QOS (quality of service) and performance requirement of different applications and services. Existing consistency definitions are tied to specific storage system implementations, and not generic enough to be used

in the context of a different (newer) storage system. It is important for consistency and isolation definitions to be implementation-independent, because a particular storage system implementation may not be suitable for different application domains. A relational storage system, like MySQL or Oracle, may be unsuitable for an application that operates in the scale of millions of concurrent executions, serving clients that are geographically distributed across different continents. In the latter case, a geo-distributed storage system like Cassandra may prove to be more efficient, producing lower latency, while compromising on serializability and strict isolation. Hence, research papers and projects often end up re-defining existing consistency models using disparate semantics and notations.

Disparities among multiple definitions of the same model may result in ambiguities and confusion, which in turn, may result in incorrect design and implementations of applications built upon such definitions. Also, with such ambiguous definitions, it is difficult to compare consistency models of different storage systems, and choose a suitable storage system for a given application. Though Adya et al. claim that their definitions are implementation agnostic, they assume that the system is able to track version histories of each object; hence, we argue that these definitions will not be applicable to systems that do not persist version histories. Further, they use graphical representation of proscribed anomalies and allowed dependencies between operations; making it difficult to automatically verify such definitions. Existing definitions of consistency models used either textual descriptions or graphical representations, making automated verification even more difficult. The goal of this paper is to specify consistency and isolation definitions using a uniform syntax in an implementation agnostic manner.

Traditionally, specifications of consistency models or isolation levels are comprised of the following components: *rules* that restrict the order in which results of operations can be observed, specified using axiomatic expressions, *anomalies* precluded by the models, expressed using directed serialization graphs (DSGs), where vertices in a DSG represent operations, and edges represent the dependency relations among operations corresponding to the connecting vertices. The rules in state-of-the-art consistency and isolation definitions specify the valid order of execution among operations in terms of dependency relations, like write-write (ww), write-read (wr) dependency, etc. Such dependency relations belong to two categories - explicit dependencies, like ww dependency, and implicit dependencies, like read-write (rw) dependency, that are expressed in terms of one or more explicit dependencies. These dependency relations are not apparent from the statement of consistency and isolation definition, especially implicit dependencies, since those are further described in terms of explicit dependencies.

However, a given consistency or isolation level specifies a valid

relative temporal order among operations. Such temporal order can be directly expressed using LTL [10] operators, like next, eventually, etc. We argue that consistency and isolation levels can be readily expressed in term of LTL-like expressions that specify required temporal order among operations according to given consistency or isolation level. Hence, we follow in the footsteps of [6] to develop a specification language *ConSpec* that is an extension of Linear Temporal Logic (LTL). LTL is a natural choice for specifying relative temporal ordering among operations in different consistency and isolation definitions. ConSpec uses LTL-like semantics to express the rules that specify the allowed order in which results of operations can be visible to a client application. The ConSpec expressions specify the allowed temporal order (i.e., “happens before” order [2]) in which operations within a client application return results, and restricts the manner in which these results are visible to subsequent operations in the client application.

We argue that specifications expressed in terms of ConSpec are much easier to reason about than the state-of-the-art definitions, which require an understanding of the various dependency relationships; the implicit dependencies can be particularly complex to express and comprehend. The main contributions of this paper are as follows.

- It presents the syntax and design of ConSpec, a language for formally specifying the consistency models and isolation levels that a storage system supports. ConSpec specifications are expressed in terms of Linear Temporal Logic (LTL), thus enabling us to leverage existing body of Satisfiability solvers and Theorem provers to build automated verifiers.
- It restates the CAP theorem, which originally used to imply linearizability when talking about consistency, taking into account the different possible consistency models. We show how the extended CAP statement segregates the lattice of consistency models into *CAP-strong* and *CAP-weak* consistency models, i.e., into consistency models which can be implemented while simultaneously providing high availability and partition tolerance, and which can not be implemented in the above manner, respectively.

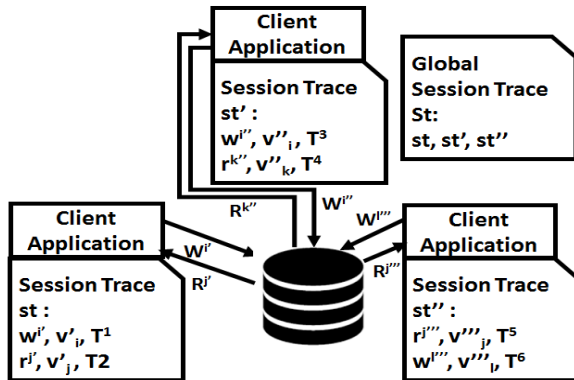


Figure 1: System Model

2. SYSTEM MODEL

We do not assume anything regarding the internal organisation of the object in the storage system. For the intent and purpose of this paper, an object o_j is a variable comprising: 1) a unique name, given by the string o_j , and 2) a value v_j , which can be of

any datatype, such as float, double, string, etc. We consider a simple execution model where client applications perform only read and write operations. More complex execution models can simply be expressed as specialised versions of basic read and write operations. A *session trace* st is a trace comprising the results of a local execution, i.e., results of operations invoked from a client application, observed from the client machine. A given client application is executed on a target storage system. In the course of execution of the client application, we collect session traces over a pre-defined sliding window over time from the client machine (i.e., the machine that hosts the client application). An operation op_j is identified by the unique signature of the operation, which is a string that formally describes the operation. The signature of a write operation op_j , that writes a value v_l to an object o , is given as $w^j(o, v_l)$, where o is a string that denotes the unique name of an object o , and v_l is a string that denotes the value that the operation writes to o . Similarly, the signature of a read operation op_k , that reads a value v_m from an object o , is $r^k(o)v_m$.

A session trace st is a collection of tuples tup_i , where each tuple is denoted as $\langle w^i, v_i', T^j \rangle$ or $\langle r^i, v_i', T^j \rangle$. The fields comprising each tuple tup_i are: 1) w^i and r^i are abbreviated signatures of write and read operations performed from a client application Cl , 2) v_i' is the value written or read by the above operation, 3) T^j is the timestamp of the operation execution, i.e., the timestamp when the operation was logged in the session trace. The *global session trace* St denotes the set of all session traces, which is equivalent to the global execution σ of Chockler et al., comprising all local executions. St is a total order comprising all session traces collected from all clients executing in a system, arranged according to their temporal order of execution. St collates all the tuples tup_i from all individual session traces, and orders the tuples according to according to descending order of the timestamp field T^j .

We do not assume anything regarding how these operations are internally handled by the storage system. For all intents and purposes, the storage system is a black box component on which read and write operations get executed. A client application executes these storage operations as high level API methods invoked from the application code. These API methods are translated into storage system specific native commands by the storage system driver which acts as an interface between the client application and the underlying storage system. Optionally, read and write operations invoked from a transaction tx are denoted as $r_{tx}^k(o)v_m$ and $w_{tx}^j(o, v_l)$, respectively. Further, we do not assume anything regarding the nature of the storage system — whether it is a standalone database, a transactional store, or a distributed storage system. Also, we do not make any assumptions regarding the configuration of the storage system, e.g., we do not assume what partitioning scheme or what normalization form (in case of RDBMS) a storage system follows. Our specifications restrict how client applications view results of storage operations, and is independent of how the underlying storage system internally processes storage operations. By not assuming anything regarding the implementation details, we render our specifications implementation independent, i.e., not specific to any particular storage system implementation. We further assume that a read operation $r_{tx}^k(o)v_m$ always returns a unique value v_m for an object o , and a write operation $w_{tx}^j(o, v_l)$ updates the value of an object o to a unique value v_l . This allows us to identify the particular write operation responsible for a particular value observed from a client application.

The specifications comprise some special purpose LTL operators. Such operators are often qualified by scope operators that specify the context within which expression is restricted. For example, we define a special-purpose operator X^o , which is derived

from the LTL modal operator “next”, i.e., X . The superscript o in the above operator is used to specify that the operator defines the LTL relationship “next” among only the operations performed specifically on the object o . The superscript o specifies the context of an object o for the enclosed relation; o is equivalent to an operand for a storage operation. An expression $W^{j'} X^o R^{k'}$ denotes that among operations performed on the object o , an execution of the operation $r^{k'}(o)v_m$ immediately follows an execution operation $w^{j'}(o, v_l)$; the superscript $'$ denotes an instance of execution of a given operation in a session trace. We also use the LTL operator “eventually”, i.e., F , to denote that an operation is preceded by the another operation in a given execution sequence. The superscript o in the above operator is used to specify that the operator defines the LTL relationship “eventually” among only the operations performed specifically on the object o .

A *serialization* is a linear sequence comprising all operations comprising a global execution on a storage system. A serialization is said to be *legal* if every read operation returns the value written by the latest write operation preceding it in the serialization. Following Chockler et al. [5], ConSpec defines consistency levels in terms of valid legal serialization orders equivalent to a given execution. Similar to Chockler et al., ConSpec defines different consistency models in terms of different specialized forms of legal serialization, which have been defined below. ConSpec defines those consistency models, which can not be expressed in terms of legal serialization, in terms of constraints on the global execution order, represented by the global session trace St . Let the propositional variable $Op^{i'}$ denote an execution of the i 'th operation comprised in the set $mathcal{O}$ of operations executed from a given client application in a session trace.

DEFINITION 1. (Legal Serialization) We represent a global execution, comprising concurrent operations executed from multiple clients on a system, with a global session trace St , which is a set of operations from all clients. A legal serialization Ser of a given global execution is represented by a sequence of propositional variables $Op^{i'}$, each variable $Op^{i'}$ denoting an execution of an operation op^i , comprised in the set of operations \mathcal{O} executed by a given client, observed a global session trace St , as follows. $Ser = \{Op^{i'} \mid Op^{i'} \in st\}$, such that $\forall op^{i'}, op^{j'} \in \mathcal{O}$
 $(Op^{i'} = W^{i'} \wedge Op^{j'} = R^{j'} \wedge \exists st \in St ((Op^{i'} F Op^{j'} \in st \wedge \neg Op^{k'} = W^{k'} (Op^{i'} F Op^{k'} F Op^{j'} \in st)) \rightarrow (v_j' = v_i'))$.

DEFINITION 2. (Serialization of a Client) We represent a local execution of a given client observed at a client machine, with a local session trace st , comprising results observed for operations executed by a given client application, comprised in the set of operations \mathcal{O} . We formally define serialization of a client, S_c , as the restriction of a legal serialization of a global execution to the operations invoked by a particular client application, which in turn, corresponds to operations observed in a particular session trace st . $S_c = \{op^{i'}, op^{j'} \in \mathcal{O} \mid Op^{i'} \mid Op^{j'} \in Ser \wedge Op^{i'} \in st\}$.

DEFINITION 3. (Serialization of a Partial Execution) The notation S_p is used to denote a legal serialization of a partial execution comprising all operations from the set of operations \mathcal{O} performed by a given client application, represented by a local session trace st , and all write operations executed from all other concurrent client applications, comprised in the global session trace St . The above partial execution is denoted as $\sigma|i + w$ by Chockler et

al. Thus, we formally denote S_p as

$$S_p = \{Op^{i'} \mid (Op^{i'} \in st) \vee (Op^{i'} = W^{i'} \wedge Op^{i'} \notin st)\}, \text{ such that } \forall op^{i'}, op^{j'} \in \mathcal{O} (Op^{i'} = W^{i'} \wedge Op^{j'} = R^{j'} \wedge \exists st \in St ((Op^{i'} F Op^{j'} \in st \wedge \neg Op^{k'} = W^{k'} (Op^{i'} F Op^{k'} F Op^{j'} \in st)) \rightarrow (v_j' = v_i'))).$$

The above legal serialization does not comprise read operations from other clients, hence its a serialization of a partial execution.

3. SYNTAX OF CONSPEC

Using the ConSpec modelling language, we develop specifications for commonly used consistency models and standardised isolation levels. The ConSpec specification for each consistency model imposes restrictive conditions on the precedence or results observed for operations performed from client applications on the underlying storage system. Such a restrictive condition is expressed as an LTL expression E^s , which is a conjunction of constraints, composed of conditions that restrict the precedence or observed results for operations performed on the system. The generalized syntax of the specification E^s for a given consistency model or an isolation level is defined as follows.

DEFINITION 4. Generalized form of ConSpec: Given a global session trace St representing a global execution of all concurrent clients on a given storage system, we say that St satisfies a consistency model \mathcal{C} , if there exists a partial order (St, \preceq) over the elements of all sequences (i.e., all local session traces) in St , such that 1) every operation op^i in St produces an output that corresponds to the result of executing sequentially a linear extension of the operations preceding op^i in \preceq , and 2) (St, \preceq) obeys E^s , which is an LTL expression restricting (St, \preceq) .

The condition 1 for the Definition 4 follows directly from the definition of St . Each element in the set St represents the event that the result of execution of an operation is observed from a client. Thus, there exists a temporal relation between at least two elements in St , denoting the temporal order among the operations, represented by a temporal operator R (which may be the temporal operator eventually F , the special-purpose next operator X^o , or a propositional relation between values returned by the operations). Condition 2 can be expressed as $E^s \models (St, \preceq)$, where E^s is the ConSpec specification for the consistency model \mathcal{C} , \models is the satisfies operator, respectively. This implies that the consistency specification for a particular consistency model \mathcal{C} must satisfy the conditions of the generalized partial order relation, i.e., Condition 1 of Definition 4.

Chockler et al. [5] define a consistency model in terms of an equivalent legal serialization, comprising a linear arrangement of the operations performed by a client application, that the model allows. Following Chockler et al., wherever possible, we also define consistency models in terms of conditions imposed with respect to a legal serialization equivalent to a given sequence of operations executed in a session trace. By definition 2, an equivalent legal serialization is a special instance of the partial order \preceq in Definition 4, thus satisfying Conditions 1 and 2 of Definition 4. Such consistency models are specified in the form $E^s = (Op^{i'} R Op^{j'} \in st \wedge \dots \wedge Op^{k'} R Op^{l'} \in st) \rightarrow$

$\exists Ser (Op^{i'} R' Op^{j'} \in Ser)$. However, some consistency models, like Total store order, Processor consistency and all isolation levels, can not be expressed in terms of an equivalent legal serialization, because of the absence of notions of legal serialization in the original definitions of these consistency models and isolation

levels. With ConSpec, those consistency models and isolation levels are specified in terms of a conjunction of logical constraints on the given global execution comprising operations performed from concurrent clients on the system. Hence, such specifications also impose a partial order of the form Definition 4, thus satisfying Conditions 1 and 2 of Definition 4. These consistency models are specified in the form $E^s = C = Op^{i'}, Op^{j'}, Op^{k'}, Op^{l'} \in St$ $(Op^{i'} R Op^{j'} \wedge \dots \wedge Op^{k'} R' Op^{l'})$.

In the next section, we specify each consistency or isolation level in terms of an expression E^s corresponding to each consistency or isolation level. The consistency model Strict Serializability is defined in terms of LTL constraints on the order among all operations performed each client in a global execution (represented by the global session trace St) with respect to an equivalent legal serialization of a client S_c . For Strict serializability, E^s is represented by a conjunction of one or more LTL constraints on operations comprised in S_c . Since S_c comprises all operations from all clients in St , the strict serializability constraint is a *total order* relation over concurrent clients executing on the underlying storage system. In other words, strict serializability imposes a total order over a global execution comprising all clients, thus imposing a total order with respect to the global session trace St . Processor consistency and Total Store Order, which are not defined in terms of legal serialization, impose constraints on operations comprised in the global execution composed of operations performed by one or more concurrent clients. Hence, for Processor consistency and Total Store Order, the specifications E^s comprise a conjunction of LTL constraints on operations comprised in the global execution St . This, in turn, implies that all of the the above mentioned consistency models can be regarded as total orders over operations performed by all clients executing on the storage system, expressed in terms of total order relations on the global session trace St in their respective specifications E^s .

On the other side of the consistency spectrum, the consistency models, namely Read Your Write, Monotonic Read, Causal consistency, Write Follows Read, and Monotonic Write are defined in terms of constraints on an equivalent legal serialization of a partial execution S_p (refer to Definition 4.1), which comprises all operations performed by a given client, along with only write operations from other clients. Specifications of these consistency models, E^s , comprise a conjunction of one or more LTL constraints on operations comprised in S_p . Since S_p does not comprise read operations performed by other concurrent clients, consistency models which impose constraints on S_p can only be regarded as *partial orders* over a global execution comprising operations performed by all concurrent clients on the underlying storage system. An LTL formula representing the specification of such a consistency model can not be regarded as a total order over operations performed by all clients (represented by the global session trace St). By definition, an LTL formula satisfies the following properties: 1) irreflexivity, since LTL defines a linear order over time, and this linearity implies that there can not be any cyclic relation in LTL formulas, 2) transitivity, which follows trivially from the linear order imposed by LTL, and 3) asymmetric, which follows directly from conditions 1 and 2. Since the above group of consistency models are specified in terms of LTL constraints, they, in turn, are irreflexive, transitive, and asymmetric. Hence, by definition, such consistency models are regarded as *strict partial orders* over all clients.

4. LIST OF CONSPEC SPECIFICATIONS FOR VARIOUS CONSISTENCY MODELS AND ISOLATION LEVELS

Here, we present the ConSpec specifications for most common consistency models and isolation levels documented till date in the literature [5, 12, 4]. We provide the detailed derivation of these specifications later in Section 7.

4.1 Specifications for Consistency Models

The Read Your Writes (RYW) model is stated as follows. Let us consider that a client application performs a write operation followed by a read operation on an object o , observed in any given session trace, obtained by executing the above client application. According to RYW, a read operation in the session trace must not observe the result of an earlier write operation, i.e., it may not observe a value written before the occurrence of the write operation which immediately precedes the read. In other words, RYW specifies: if an operation r^j reads an object o that was written by a write w^i , the r^j must not return the value written by write operation w^k that may have preceded w^i . In other words, it requires that $v_i \neq v_k$, where \neq denotes the traditional arithmetic operator not equals. We use serialization of a partial execution S_p , specified in Definition 4.1, to denote a legal serialization of all operations comprising a given client application, as well as all write operations executed from all other concurrent client applications. Following the format of Definition 4, the above RYW definitions are reduced to LTL-like expression, as follows.

$$\mathcal{E}^s = \forall st \in St, W^{i'}, R^{j'} \in st ((W^{i'} F R^{j'} \in st) \rightarrow (W^{i'} \preceq R^{j'})), \quad (1)$$

where the partial order (St, \preceq) restricts the order among conflicts pair of writes and reads in an equivalent serialization of a partial execution S_p , defined in Definition 4.1.

Read/Session Monotonic (MR) specifies: read operations on a storage system object invoked from a client application or a session must always return results in an increasing order of recency as per their mutual order in the session trace. MR is formally represented by the following expression.

$$\mathcal{E}^s = \forall st \in St, R^{i'}, R^{j'} \in st (R^{i'} F R^{j'} \in st \rightarrow (R^{i'} \preceq R^{j'})), \quad (2)$$

where (St, \preceq) constrains the order among conflicting reads in an equivalent S_p .

Causal Consistency specifies: results of operations (in a given client application) that are bound by a causal relationship [2] must be visible in the session trace according to their mutual execution order with respect to an observed session trace. The ConSpec expression for Causal Consistency is given below.

$$E^s = \forall st \in St, Op^{i'}, Op^{j'} \in st ((Op^{i'} F Op^{j'} \in st) \rightarrow (Op^{i'} \preceq Op^{j'})), \quad (3)$$

where the partial order (St, \preceq) enforces order among conflicting operations in an equivalent S_p to follow the causal order.

Strict Serializability specifies: any session trace must observe latest result of all operations in a client application. Using the notion of serialization of a client, given in Equation 2, Strict Serializability is expressed as follows.

$$E^s = \forall st \in St, Op^{i'} \in st, Op^{j'} \in st (Op^i X^o Op^{j'} \in st \rightarrow (Op^{i'} \preceq Op^{j'})), \quad (4)$$

where the partial order (St, \preceq) restricts an equivalent serialization of a client S_c , defined in Definition 2, to respect the order of successive operations in a given session trace.

Write Follows Read (WFR) specifies: a write operation, that follows a read operation on a common storage system object, must happen while the object contains a value that is at least as old as the one that was read by the preceding read operation. WFR is expressed as follow.

$$E^s = \forall st \in St, R^{i'}, W^{j'} \in st ((R^{i'} FW^{j'} \in st) \rightarrow (R^{i'} \preceq W^{j'})), \quad (5)$$

where the partial order (St, \preceq) constrains the order among conflicting pairs of reads and writes in an equivalent S_p .

Monotonic Writes (MW) model is specified as: successive write operations invoked from a client application on a storage system object must occur according to their mutual order in the session trace. MW is expressed as follows, in terms of a legal serialization for the given execution, specified in Definition .

$$E^s = \forall st \in St, W^{i'}, W^{j'} \in st (W^{i'} FW^{j'} \in st \rightarrow (W^{i'} \preceq W^{j'})), \quad (6)$$

where the partial order (St, \preceq) restricts the order among conflicting writes in an equivalent S_p .

Owen et al. [11] specifies Total Store Order (TSO) as: read operations performed on the same storage system object by client applications can not return a subsequent (or newer) value unless the result of a given write is observed by all client applications, i.e., observed in all session traces. A session trace for a client application can not view results of a sequence of write operations in an order which differs from the invocation order. Following Definition 4, TSO is expressed as follows.

$$E^s = \forall x, st \in St \neg \exists op^i(x), op^j(x) \in \mathcal{O} (Op^{i'}(x) F Op^{j'}(x) \in st) \rightarrow (Op^{j''}(x) F Op^{i''}(x) \in st),$$

which can be rewritten as follows.

$$\forall x, (Op^{i'}(x) F Op^{j'}(x) \in st) \rightarrow (Op^{i'}(x) \preceq (x) Op^{j'}(x)), \quad (7)$$

where $\preceq(x)$ is a partial order over a given object x with respect to a given global session trace St .

Processor Consistency specifies: operations in a client application must execute according to the invocation order with respect to a given processor (client); i.e., result of a preceding operation invoked from a given client must be observed in all processors before result of a succeeding operation can be observed in any processor (or client application). Reads against write operations should view the results of the writes according to the invocation order of these writes in a client application. Thus, Processor Consistency can be expressed as

$$E^s = \forall x, y, st, st' \in St \neg \exists op^i(x), op^j(y) \in \mathcal{O} (Op^{i'}(x) F Op^{j'}(y) \in st) \rightarrow (Op^{j''}(y) F Op^{i''}(x) \in st'),$$

which can be rewritten as follows.

$$\forall x, y, st \in St (Op^{i'}(x) F Op^{j'}(y) \in st) \rightarrow (Op^{i'}(x) \preceq Op^{j'}(y)), \quad (8)$$

where \preceq is a partial order over operations across all data objects with respect to a given global session trace St .

4.2 Specification of Isolation Levels

For specifying isolation levels, we consider only the isolation levels defined by Adya et al. [1]. We show the equivalence between the graphical representation of the dependency relations used by Adya et al. and the LTL specifications of ConSpec in Section 9. The isolation level PL-1 mandates that write operations in a transaction must be executed such that write-write dependency cycles among transactions are proscribed. Following Definition 4, PL-1 is specified as follow.

$$E^s = \forall st \in St, tx, ty \in St, x, y, W^i(x)'_{tx}, W^j(y)'_{tx}, W^k(x)'_{ty}, W^l(y)'_{ty} \in St ((\neg R^m(x)'_{tx}, R^n(y)'_{ty} \in St) ((v'_n = v'_l) \wedge (v'_m = v'_i)) \vee ((v'_n = v'_j) \wedge (v'_m = v'_k))) \quad (9)$$

where tx and ty are consecutive transactions.

The isolation level PL-2 states that operations from a transaction can not read a value that was not the final value that was committed by a transaction. Further, a transaction cannot read values written by aborted transactions. Additionally, operations in a transaction must be executed in an order that proscribes write-write and write-read dependency cycles [1] among transactions. PL-2 is given as

$$E^r = \forall st, st' \in St, tx, ty \in St (((\neg \exists a_{tx} \in st (W^i(x)'_{tx} F R^j(x)'_{ty} \in st) \wedge \neg \exists a_{ty} \in st (W^k(y)'_{ty} F R^l(y)'_{tx} \in st)) \wedge (\neg \exists W^i(x)'_{tx}, W^j(x)'_{tx}, R^k(x)'_{ty} \in st) ((W^i(x)'_{tx} F W^j(x)'_{tx} F c_{tx} \in st) \wedge (v'_k = v'_i)) \wedge (\neg (W^i(x)'_{ty} F W^j(x)'_{ty} F c_{ty} \in st) \wedge (v'_k = v'_i)) \wedge (\neg \exists W^i(x)'_{tx}, W^j(y)'_{tx}, W^k(x)'_{ty}, W^l(y)'_{ty}, R^m(x)'_{tx}, R^n(y)'_{ty} \in St) ((v'_n = v'_l) \wedge (v'_m = v'_i)) \vee ((v'_n = v'_j) \wedge (v'_m = v'_k)))) \quad (10)$$

where c_{tx} and c_{ty} are the commit statements for the transactions tx and ty , respectively. F is the traditional LTL next operator.

The isolation level PL-3 specifies that, on top of PL-2, read-write dependency cycles [1] must also be proscribed among transactions.

PL-3 is given as

$$\begin{aligned}
E^r = C = \forall st, st' \in St, tx, ty \in St \\
& (((\neg a_{tx} \in st(W^i(x)_{tx}' FR^j(x)_{ty}' \in st) \\
& \wedge \neg a_{ty} \in st(W^k(y)_{ty}' FR^l(y)_{tx}' \in st)) \wedge (\neg W^i(x)_{tx}', W^j(x)_{tx}', \\
& R^k(x)_{ty}' \in st) \wedge (W^i(x)_{tx}' FW^j(x)_{tx}' FC_{tx} \in st) \\
& \wedge (v_k' = v_i')) \wedge \neg (W^i(x)_{ty}', W^j(x)_{ty}', R^k(x)_{tx}' \in st) \\
& ((W^i(x)_{ty}' FW^j(x)_{ty}' FC_{ty} \in st) \wedge (v_k' = v_i')) \wedge \\
& \neg W^i(x)_{tx}', W^j(y)_{tx}', W^k(x)_{ty}', W^l(y)_{ty}', R^m(x)_{tx}', R^n(y)_{ty}' \in St \\
&)) \wedge ((v_n' = v_l') \wedge (v_m' = v_i')) \vee ((v_n' = v_j') \wedge (v_m' = v_k')) \wedge \\
& (\neg R^i(x)_{tx}', W^j(x)_{ty}', R^k(x)_{ty}', R^l(y)_{ty}', W^m(y)_{tx}', R^n(y)_{tx}' \in St) \\
& ((R^i(x)_{tx}' FW^j(x)_{ty}' FR^k(x)_{ty}' \in St) \wedge \\
& (R^l(y)_{ty}' FW^m(y)_{tx}' FR^n(y)_{tx}' \in St)) \\
& ((v_k' = v_j') \wedge (v_n' = v_m')))) \quad (11)
\end{aligned}$$

The isolation level PL-2.99 specifies that, on top of PL-2, read-write item-dependency cycles must be proscribed among transactions. PL-2.99 is given as

$$\begin{aligned}
E^r = C = \forall st, st' \in St, tx, ty \in St \\
& (((\neg a_{tx} \in st(W^i(x)_{tx}' FR^j(x)_{ty}' \in st) \\
& \wedge \neg a_{ty} \in st(W^k(y)_{ty}' FR^l(y)_{tx}' \in st)) \wedge (\neg W^i(x)_{tx}', W^j(x)_{tx}', \\
& R^k(x)_{ty}' \in st) \wedge (W^i(x)_{tx}' FW^j(x)_{tx}' FC_{tx} \in st) \\
& \wedge (v_k' = v_i')) \wedge \neg (W^i(x)_{ty}', W^j(x)_{ty}', R^k(x)_{tx}' \in st) \\
& ((W^i(x)_{ty}' FW^j(x)_{ty}' FC_{ty} \in st) \wedge (v_k' = v_i')) \wedge \\
& \neg W^i(x)_{tx}', W^j(y)_{tx}', W^k(x)_{ty}', W^l(y)_{ty}', R^m(x)_{tx}', R^n(y)_{ty}' \in St \\
&)) \wedge ((v_n' = v_l') \wedge (v_m' = v_i')) \vee ((v_n' = v_j') \wedge (v_m' = v_k')) \wedge \\
& (\neg R^i(x)_{tx}', W^j(x)_{ty}', R^k(x)_{ty}', R^l(y)_{ty}', W^m(x)_{tx}', R^n(x)_{tx}' \in St) \\
& ((R^i(x)_{tx}' FW^j(x)_{ty}' FR^k(x)_{ty}' \in St) \wedge \\
& (R^l(y)_{ty}' FW^m(x)_{tx}' FR^n(x)_{tx}' \in St)) \\
& ((v_k' = v_j') \wedge (v_n' = v_m')))) \quad (12)
\end{aligned}$$

5. REWRITING CAP THEOREM IN TERMS OF CONSPEC

CAP theorem [3, 7] states that it is impossible to simultaneously achieve atomic consistency (linearizability), availability (i.e., the property that there must be a response against every request), and partition tolerance (i.e., tolerance of network partitions) in a distributed system. However, the original statement of CAP is restricted to linearizability, and does not address the wide array of consistency models supported by modern storage systems. Indeed modern storage systems relax consistency guarantees to favour availability and partition tolerance according to the use case and system requirements. The various degrees of relaxation gives rise to a lattice of consistency models, ranging from RYW to Strict Se-

rializability. In this section, we rewrite the CAP theorem to take into account this wide array of consistency models. Here, we analyze the specifications of consistency models, and derive the condition that makes a particular consistency model CAP-strong, i.e., strong enough to be bounded by the constraints of CAP theorem, as follows. First, we state and prove the condition for a consistency model to be CAP-strong, and then we state and prove the opposite case, i.e., we prove the condition for a consistency model to be CAP-weak.

THEOREM 1 (EXTENDED CAP THEOREM). *It is impossible for a storage system to satisfy a consistency model that enforces restrictions on order among the results of operations performed by different clients executing on the system, while simultaneously providing high availability and partition tolerance.*

To elaborate, if the specification E^s (following the generic specification format given in Definition 4) of a consistency model imposes a partial order \preceq among results observed by pairs of operations performed by different clients executing on underlying storage system (i.e., among operations from different clients comprised in a global execution represented by a global session trace St), that consistency model is regarded as a CAP-strong consistency model. On the other hand, a consistency model that imposes a partial order \preceq only among results observed by pairs of operations performed by the same client comprised in a global execution represented by a global session trace St , is regarded as CAP-weak. It directly follows from Theorem 5 that the CAP-weak condition, in turn, implies that the specified partial order \preceq as per Definition 4 must be a partial order on only operations from the same client, comprised in St . In other words, for a CAP-weak consistency model, Condition 1 in Definition 4 can be rewritten as the following condition. Thus, the condition $\forall Op^i, Op^j \in St : Op^i \preceq Op^j \vee Op^j \preceq Op^i$ must hold for operations $op^i \in st$ and $op^j \in st'$ comprised in separate session traces st and st' comprising results observed by operations performed by different clients Cl_1 and Cl_2 .

Proof: Following the proof style of Glibert et al. [7], we prove the above theorem by contradiction as follows. Let us consider a pair of concurrent clients, namely Cl_1 and Cl_2 executing on two storage server nodes N_1 and N_2 . Let us consider that the network between N_1 and N_2 is partitioned at a particular instant T_i of execution of the clients, i.e., all communication between Cl_1 and Cl_2 are lost. Further let us assume that Cl_1 and Cl_2 perform operations on a common data object x . The clients perform a sequence of write operations α onto x until a network partition occurs, and they are placed in two disjoint components of the network G_1, G_2 . Since the system is still available, both servers accept client requests, and apply the corresponding writes onto local copies of x . Because of this, Cl_1 (located in N_1) applies a new sequence of writes β onto x while Cl_2 (located in N_2) applies a different sequence of writes γ to x . Since these components are disconnected, none of these writes are propagated among the servers N_1 and N_2 . Therefore, the global session trace St , as observed by the client Cl_1 , is a sequence α, β for the conflicting object x ; whereas St , as observed by Cl_2 , is a sequence α, γ . Let operations op_i and op_j be the prefix of the executions β , and γ observed by clients Cl_1 and Cl_2 , respectively. Consider that the storage system enforces a consistency model which imposes a partial order \preceq among results of operations from different concurrent clients executing on the underlying storage system. Since \preceq enforces a partial order among results of operations performed by different concurrent clients, the results observed by clients Cl_1 and Cl_2 must be restricted to a common partial order \preceq . This implies that results of operations op_i and

op_j in each of the global session traces observed by the clients Cl_1 and Cl_2 must either satisfy $op_i \preceq op_j$ or $op_j \preceq op_i$, respectively. Since $\beta \neq \gamma$ in St , hence $op_i \preceq op_j$ holds in β , while $op_j \preceq op_i$ holds in γ ; hence the above condition of partial order among operations from different clients is violated in the global session traces α, β and α, γ . This contradicts the requirement of a common partial order \preceq on the operations performed by different clients comprised in St . Hence, such a consistency model can not achieve high availability and partition tolerance at the same time.

THEOREM 2. *It is possible for a storage system to satisfy a consistency model that enforces a partial order among results observed by operations performed by the same client executing on the system, while simultaneously providing high availability and partition tolerance.*

Following the approach of the proof to Theorem 5, the above theorem is proved as follows. We consider two servers N_1 and N_2 that serve requests from clients Cl_1 and Cl_2 . We consider that a partition occurs at a particular instant between the servers N_1 and N_2 . Let us consider that the underlying storage system satisfies a consistency model \mathcal{C} that imposes a partial order \preceq on a global execution comprising operations performed by all clients (represented by a global session trace St). The global session traces observed by clients Cl_1 and Cl_2 comprise sequences of operations, given as α, β and α, γ , respectively. Because of the lack of communication of updates between the servers N_1 and N_2 serving Cl_1 and Cl_2 in presence of a partition, the global session trace observed by Cl_1 satisfies the partial order $\preceq = \alpha \preceq \beta$ or $\beta \preceq \alpha$; whereas the global session trace observed by Cl_2 satisfies a different partial order $\preceq' = \alpha \preceq \gamma$ or $\gamma \preceq \alpha$. Since $\beta \neq \gamma$, each of the above session traces satisfies a partial order \preceq and \preceq' . Hence, in presence of partitions, the above system remains available, while satisfying a partial order on the global session trace St . Thus, such a consistency model can be implemented in a manner which simultaneously achieves high availability and partition tolerance.

6. ANALYSIS OF CONSISTENCY MODELS WITH RESPECT TO CAP

For RYW, MR, Causal consistency, WFR, and MW, E^s imposes restrictive conditions on operations comprised in an equivalent legal serialization of a partial execution S_p , composed of operations performed by the given client, along with only writes from other concurrent clients. Since S_p does not comprise read operations from all clients in St , the above consistency models impose only a strict partial order with respect to operations in St . Thus, writes from concurrent clients are allowed to overwrite results of writes from a given client as long as there exists an equivalent legal serialization comprising the combination of operations from the given client and the concurrent writes. It only restricts that the partial order, i.e., precedence, among the operations from each client is maintained. Similarly, Processor consistency (PC) restricts that the order among operations in a given client must be preserved in all executions, i.e., no client may observe the operations from a given client in a different order than the invocation order. Hence, PC does not enforce a partial order across operations from different concurrent clients, it only imposes a strict partial order on all operations from a given client. Thus, these comprise a weaker class of consistency models that can be satisfied without synchronization among clients, i.e., these consistency models are not dependent among reliable communication among concurrent clients (like broadcasting of results on common data objects, or blocking of conflicting operations). Using the expression for S_c from Definition 2

in Equation 4, the specification for Strict Serializability reduces to

$$E^s = \forall i, j, st (Op^i X^o Op^{j'} \in st \rightarrow \exists (Op^{i'}, Op^{j'} \in Ser \wedge Op^{i'}, Op^{j'} \in st) (Op^{i'} X^o Op^{j'} \in S_c)).$$

Since, $S_c \subset Ser$, $Op^{i'} X^o Op^{j'} \in S_c$ implies $Op^{i'} X^o Op^{j'} \in Ser$. Hence, moving the existential quantifier out of the enclosing scope of the universal quantifier, the above expression can be rewritten as

$$E^s = \exists (Op^{i'}, Op^{j'} \in Ser) \forall i, j, st (Op^i X^o Op^{j'} \in st \rightarrow (Op^{i'} X^o Op^{j'} \in Ser)).$$

Thus, Strict Serializability restricts that the mutual execution order among operations performed by each client in a global execution must be preserved in an equivalent legal serialization Ser , comprised of all operations in a global execution. Thus, Strict serializability imposes a partial order among operations from all clients comprised in St . TSO imposes restrictions on the execution order and observed results among all operations performed by different concurrent clients observed in a global session trace St . Hence, it directly follows from the definitions that TSO imposes a partial order on St , where each operation performed by a given client in St is bounded by a partial order relation with an operation in St performed by another concurrent client. Hence, these consistency models require strict synchronization among operations from different clients, such that operations from each client must be executed in isolation from other concurrent clients.

7. DERIVING CONSPEC SPECIFICATIONS FROM THE STATE-OF-THE-ART DEFINITIONS

Here we show how the ConSpec specifications can be directly derived from state-of-the-art specifications of consistency and isolation level [5, 12]. As already discussed, state-of-the-art consistency specifications define axiomatic rules for consistency models using first order logic expressions. Consistency definitions are specified by Chockler et al. [5] as follows. They use the following formalization syntax in their definitions. Consider two operations op^1 and op^2 invoked from a client application, such that op^1 is a write operation that updates a storage system object o , and op^2 is a read operation that reads o . The symbol \rightarrow denotes precedence relationship [2] between two operations. Chockler et al. used the notation σ to denote an execution of a sequence of read and write operations performed by a client application (referred to as a process by Chockler). They also use a specialised form of the precedence operator, $\xrightarrow{\sigma}$, where the superscript σ is used to indicate that the precedence relationship is defined specifically with respect to an execution σ by a client application. Thus, an expression $op^1 \xrightarrow{\sigma} op^2$ indicates that an execution of the operation op^1 precedes, i.e., happens before, an execution of the operation op^2 in an execution sequence σ . Chockler et al. denote a partial execution comprising all operations performed by a process (or a client application) and all write operations from other processes as $\sigma[i + w]$. They denote a legal serialization for the above partial execution as S_i ; we denote the above legal serialization as S_p . For means and purposes of this paper, an execution sequence σ according to the terminology of Chockler et al. is equivalent to a session trace st in ConSpec. Hence, we assign the notation st to an execution sequence σ in Chockler's definitions. We can rewrite the precedence relation $op^1 \xrightarrow{\sigma} op^2$, in terms of Linear Temporal Logic (i.e., LTL) $op^1 F op^2 \in st$, where we replace the prece-

dence operator \rightarrow with the equivalent LTL operator F that denotes “eventually”. The definitions by Chockler et al. express a precedence relation $op^1 \rightarrow op^2$ in a serialization S_i as $op^1 \xrightarrow{S_i} op^2$, where we specify the scope of the precedence operator \rightarrow with the serialization S_i . We can rewrite the above precedence relation as $op^1 \xrightarrow{S_i} op^2$, in terms of Linear Temporal Logic (i.e., LTL) $op^1 F op^2 \in S_i$, where we assign the notation S_i to a serialization S_i , and replace the precedence operator with the equivalent LTL operator F that denotes “eventually”. Also, Chockler et al. bases their definitions on a serialization of an execution sequence comprising all operations performed by a client application (or a process), along with writes by all other clients. In ConSpec, we denote such a serialization as S_p . We can express the precedence relation for such a serialization in terms of LTL as $op^1 F op^2 \in S_p$.

Let w^i and r^j be signatures of read and write operations op^1 and op^2 , respectively. Chockler et al. states the RYW consistency model as: if the condition represented by the expression $op^1 \xrightarrow{\sigma} op^2$ holds for a given execution sequence σ (let us call this Condition 1) comprising the above write and read operations, there must exist a serialization S_i , comprising the operations op^1 and op^2 , for which the condition given by the expression $op^1 \xrightarrow{S_i} op^2$ must hold (let us call this Condition 2). The above precedence relationships among operations op^1 and op^2 in Condition 1 can be directly expressed in terms of LTL as follows. Let Op^1 and Op^2 be propositional logic variables that indicate whether operations op^1 and op^2 on an object o have returned (indicated by the value of the variables being TRUE) or not (indicated by the value of the variables being FALSE), respectively. Condition 1, which expresses the precondition for RYW with respect to a given execution order σ , can be expressed in terms of the LTL expression $Op^1 F Op^2 \in st$, where a session trace st is equivalent to σ for the means and purposes of this paper. The above LTL expression implies: if the propositional variable Op^1 is true at one instant, Op^2 is eventually true in some later instant for the same session trace st . In other words, the operation op^2 on object o follows the operation op^1 in the given session trace st . The expression $op^1 \xrightarrow{S_i} op^2$ in Condition 2 can be expressed in terms of LTL as $Op^1 F Op^2 \in S_p$, where S_p is equivalent to S_i , i.e., both notations represent a legal serialization comprising all operations performed from a given client application along with writes from all other clients. Further, since RYW talks about a sequence of write operations followed by a read, the propositional variables Op^1 and Op^2 in both Condition 1 and Condition 2 can safely be replaced by a new propositional variable W^i and R^j . Since $W^i F R^j \in S_p$ denotes a partial order among W^i and R^j , we can rewrite the above expression as $W^i \preceq R^j$. Thus, Chockler’s definition reduces into the specification given in Equation 4.1.

Session Monotonic or Monotonic Read (also referred to as Session Causality or MR) consistency model is another popular consistency model [5, 12]. According to Chockler et al., MR is expressed as: if the condition $op^1 \xrightarrow{\sigma} op^2$ holds for a given execution sequence σ (Condition 1), where both op^1 and op^2 must be read operations, there must exist a serialization S_i , comprising the operations op^1 and op^2 , for which the condition $op^1 \xrightarrow{S_i} op^2$ must hold (Condition 2). As in the case of RYW, the expressions $op^1 \xrightarrow{\sigma} op^2$, and $op^1 \xrightarrow{S_i} op^2$ can be rewritten as LTL expressions $R^i F R^j \in st$ and $R^i F R^j \in S_p$, respectively. Again, since $R^i F R^j \in S_p$ denotes a partial order among R^i and R^j , we can rewrite the above expression as $R^i \preceq R^j$, thus reducing the above specification into Equation 2. Thus, the ConSpec specification for MR is directly derived from the definition by Chockler et al.

Causal consistency is specified by Chockler et al. as follows. Chockler et al. defines a *direct precedence relation* (which we denote as $\xrightarrow{\sigma}$) between operations op^i and op^k in the execution order σ as follows. $op^i \xrightarrow{\sigma} op^j$ implies that either op^j reads values written by operation op^i , or a precedence relationship exists in the execution order σ between op^i and op^j , i.e., $op^i \xrightarrow{\sigma} op^j$ exists. The former part of the precondition can be expressed in the form $Op^{i'} F Op^{j'}$, where the LTL-like operator F is equivalent to the precedence operator. The latter part of the precondition implies read operation Op^i reads the value written by Op^j , i.e., $(Op^{i'} = W^{i'}) \wedge (Op^{j'} = R^{j'}) \wedge (v_i = v_j)$. The precondition of causal consistency specifies that a transitive closure must exist for a direct precedence relation $\xrightarrow{\sigma}$, i.e., $\xrightarrow{\sigma}^*$ must exist, i.e., the condition $Op^{i'} F Op^{j'} \vee ((Op^{i'} = W^{i'}) \wedge (Op^{j'} = R^{j'}) \wedge (v_i = v_j))$ must hold. Chockler et al. specifies that for causal consistency, if the above precondition holds, the condition $op^i \xrightarrow{\sigma} op^j$ must hold, i.e., an occurrence of op^i must be followed by an occurrence of op^j in any legal serialization S_i . The above postcondition can be expressed in terms of LTL as $\exists S_p (\forall Op^i F Op^j \in S_p)$. The operation op^i precedes op^j such that the condition $Op^{i'} F Op^{j'} \in st$ is satisfied. Since $Op^{i'} F Op^{j'} \in S_p$ denotes a partial order among $Op^{i'}$ and $Op^{j'}$, we can rewrite the above expression as $Op^{i'} \preceq Op^{j'}$. Thus, Chockler’s definition reduces into the specification given in Equation 3.

Chockler et al. states strict serializability as: the order of execution among operations comprised in any execution sequence of a given client application must match the precedence order in a legal serialization for the given execution sequence. The order of execution among operations in an execution sequence σ can be given in terms of a series of LTL expression that express the precedence of successive operations in the sequence. The precedence relation between any two successive operations op^i and op^j can be expressed as $op^i X op^j$. Thus, the above condition can be expressed as: any two operations in an execution sequence must execute in an order that matches the precedence relation among these operations in a legal serialization of the above execution sequence. In other words, any two operations must occur in a session trace in an order that matches the precedence relation of these operations in a legal serialization of those operations. Hence, this condition can be expressed in terms of LTL as $Op^i X Op^j \in st \rightarrow \exists S_c (Op^i X Op^j \in S_c)$. Since $Op^i X Op^j \in S_c$ denotes a partial order among Op^i and Op^j , we can rewrite the above expression as $Op^i \preceq Op^j$. Thus, Chockler’s definition reduces into the specification given in Equation 4.

Chockler et al. states the WFR consistency model as: if the condition represented by the expression $op^1 \xrightarrow{\sigma} op^2$ holds for a given execution sequence σ of read operation followed by write operation (let us call this Condition 1), there must exist a serialization S_i , comprising the operations op^1 and op^2 , for which the condition given by the expression $op^1 \xrightarrow{S_i} op^2$ must hold (let us call this Condition 2). The above precedence relationships among operations op^1 and op^2 in Condition 1 can be directly expressed in terms of an LTL expression $R^i F W^j \in st$, where a session trace st is equivalent to σ for the means and purposes of this paper. The expression $op^1 \xrightarrow{S_i} op^2$ in Condition 2 can be expressed in terms of LTL as $R^i F W^j \in S_p$, where S_p is equivalent to S_i . Since $R^i F W^j \in S_p$ denotes a partial order among R^i and W^j , we can rewrite the above expression as $R^i \preceq W^j$. Thus, Chockler’s definition reduces into the specification given in Equation 5.

Chockler et al. states the MW consistency model as: if the condition represented by the expression $op^1 \xrightarrow{\sigma} op^2$ holds for a given execution sequence σ of write operations (let us call this Condition 1), there must exist a serialization S_i , comprising the operations op^1 and op^2 , for which the condition given by the expression $op^1 \xrightarrow{S_i} op^2$ must hold (let us call this Condition 2). The above precedence relationships among operations op^1 and op^2 in Condition 1 can be directly expressed in terms of an LTL expression $Op^1 F Op^2 \in st$, where a session trace st is equivalent to σ for the means and purposes of this paper. The expression $op^1 \xrightarrow{S_i} op^2$ in Condition 2 can be expressed in terms of LTL as $Op^1 F Op^2 \in S_p$, where S_p is equivalent to S_i . Further, since MW talks about a sequence of write operations, the propositional variable Op^1 and Op^2 in both Condition 1 and Condition 2 can safely be replaced by new propositional variables W^i and W^j . Since $W^i F W^j \in S_p$ denotes a partial order among W^i and W^j , we can rewrite the above expression as $W^i \preceq W^j$. Thus, Chockler's definition reduces into the specification given in Equation 6.

Total Store Order (TSO or Total Order) is specified by Owen et al. [11] as follows. In any two executions of a sequence of operations by two processors (or client applications), a pair of operations on a common storage system object x must be executed in the same precedence order. The above condition implies that any session trace st must comprise write-read sequences $op^i(x), op^j(x)$ on an object x in the temporal order specified in the global session trace St . Thus, $E^s = \forall x, st \in St \ \neg Op^i(x), op^j(x) \in \mathcal{O}$ $(Op^i(x) F Op^j(x) \in st) \rightarrow (Op^{j''}(x) F Op^{i''}(x) \in st)$ The above expression can be further rewritten in terms of a partial order $\preceq(x)$ over a given object x as $\forall x, (Op^{i'}(x) F Op^{j'}(x) \in st) \rightarrow (Op^{i'}(x) \preceq(x) Op^{j'}(x))$, where $\preceq(x)$ is a partial order over a given object x with respect to a given global session trace St . Similarly, the specification for Processor consistency (PC) model can be directly derived using the above line of reasoning; PC talks about preserving the order of operations executed by individual processors (or client applications). Processor Consistency specifies: operations in a client application must execute according to the invocation order with respect to a given processor (client); i.e., result of a preceding operation invoked from a given client must be observed in all processors before result of a succeeding operation can be observed in any processor (or client application). Reads against write operations should view the results of the writes according to the invocation order of these writes in a client application. Consider that a given client observes operations $op^i(x), op^j(y)$ in a temporal order $Op^{i'}(x) F Op^{j'}(y)$ observed in a session trace st . Then, no other session trace st' observed by any other concurrent client may observe the operations in an order which violates the temporal order in st . Thus, Processor Consistency can be expressed as $E^s = \forall x, y, st, st' \in St \ \neg Op^i(x), op^j(y) \in \mathcal{O}$ $(Op^{i'}(x) F Op^{j'}(y) \in st) \rightarrow (Op^{j''}(y) F Op^{i''}(x) \in st')$, which can be rewritten as follows. $\forall x, y, st \in St$ $(Op^{i'}(x) F Op^{j'}(y) \in st) \rightarrow (Op^{i'}(x) \preceq Op^{j'}(y))$, where \preceq is a partial order over operations across all data objects with respect to a given global session trace St .

The ConSpec specifications for the isolation levels are derived from the definitions by Adya et al. [1]. Here, we show that the isolation specifications in ConSpec, given in Equations 9, 10, 11, and 12, follow directly from the definitions of isolation levels PL-1, PL-

2, PL-3, and PL-2.99 [1]. We analyze the ConSpec Equations, and demonstrate their equivalence to the definitions of [1], as follows.

The PL-1 specification states that the anomaly G0 must be proscribed in any session trace collected for the execution of a client application. G0 specifies that there can not be any directed cycle in the dependency graph corresponding to an execution of a pair of transactions, comprising entirely of write dependency edges. Let us consider a pair of transactions tx and ty . Let us consider that write operations $w^i(x)_{tx}$ and $w^j(y)_{tx}$ write to objects x and y from transaction tx , and $w^k(x)_{ty}$ and $w^l(y)_{ty}$ write to objects x and y from transaction ty . According to PL-1 there can not exist a pair of read operations $r^m(x)_{tx}$ and $r^n(y)_{ty}$ invoked from transactions tx and ty such that the following conditions are simultaneously satisfied: 1) $r^m(x)_{tx}$ reads the result of the write operation $w^i(x)_{tx}$ causing a ww dependency from tx to ty , i.e., $v_m = v_i$ must not hold, and 2) $r^n(y)_{ty}$ reads the result of the write operation $w^l(y)_{ty}$ causing a ww dependency between ty to tx , i.e., $v_n = v_l$ must not hold. Simultaneous satisfaction of conditions 1 and 2 results in a cycle comprising ww dependencies between tx and ty . Thus, we can express the specification of PL-1 as $C = \forall St, st, tx, ty, x, y, W^i(x)_{tx}, W^j(y)_{tx}, W^k(x)_{ty}, W^l(y)_{ty} \in St$ $(\neg R^m(x)_{tx}, R^n(y)_{ty} \in st) \vee ((v_n = v_i) \wedge (v_m = v_l)) \vee ((v_n = v_j) \wedge (v_m = v_k))$ which is identical to Equation 9.

The isolation level PL-2 proscribes the anomalies G1-a, G1-b, and G1-c [1]. We show that the ConSpec specification for PL-2, given in Equation 10, disallows the above anomalies as follows. **G1-a:** According to G1-a, a transaction can not observe a value written by an aborted transaction. In other words, if either transaction tx or transaction ty aborts, i.e., either of the conditions $a_{tx} \in st$ or $a_{ty} \in st$ hold, then there can not exist a read operation that reads the value written by an aborted transaction, i.e., the conditions $W^i(x)_{tx} F R^j(x)_{ty} \in st$ or $W^k(y)_{ty} F R^l(y)_{tx} \in st$ can not hold. Thus, the G1-a condition can be directly translated to the expression $\neg a_{tx} \in st \ (W^i(x)_{tx} F R^j(x)_{ty} \in st) \wedge$

$\neg a_{ty} \in st \ (W^k(y)_{ty} F R^l(y)_{tx} \in st)$ in Equation 10. **G1-b:** According to G1-b, a transaction must always read the final committed version of an object. Let us consider consecutive write operations $w^i(x)_{tx}$ and $w^j(x)_{tx}$ invoked from transaction tx followed by the commit statement c_{tx} , and a read operation $r^k(x)_{ty}$ invoked from transaction ty , i.e., $W^i(x)_{tx} F W^j(x)_{tx} F c_{tx} F R^k(x)_{ty} \in st$. According to G1-b, ty can never read a value written by an operation in tx that was not finally committed by tx . Thus, $r^k(x)_{ty}$ cannot return the result of $w^i(x)_{tx}$ since it was overwritten by another write $w^j(x)_{tx}$ before the commit c_{tx} . The above condition can be expressed as $\neg \exists ((W^i(x)_{tx} F W^j(x)_{tx} F c_{tx} F R^k(x)_{ty} \in st) \wedge (v_k = v_i)) \wedge \neg \exists ((W^i(x)_{ty} F W^j(x)_{ty} F c_{ty} F R^k(x)_{tx} \in st) \wedge (v_k = v_i))$. **G1-c:** G1-c specifies that there can be no direct cycle comprising dependency edges, i.e., wr and ww dependencies, between transactions in a given execution. The condition $\neg (W^i(x)_{tx}, R^j(y)_{tx}, W^k(y)_{ty}, R^l(x)_{ty} \in St) \vee ((v_l = v_i) \wedge (v_k = v_j))$ proscribes wr dependency cycles. The expression $(\neg W^i(x)_{tx}, W^j(y)_{tx}, W^k(x)_{ty}, W^l(y)_{ty}, R^m(x)_{tx}, R^n(y)_{ty} \in St) \vee ((v_n = v_l) \wedge (v_m = v_i)) \vee ((v_n = v_j) \wedge (v_m = v_k))$ proscribes ww cycles. Thus, the ConSpec specification of PL-2 in Equation 10 proscribes G1-a, G1-b, and G1-c; thus, Equation 10 is equivalent to PL-2 specifications by Adya et al.

The isolation level PL-3 proscribes the anomaly G-2 [1]. According to the ConSpec specification for PL-3 (refer to Equation 11), apart from proscribing G1 anomalies, G2 anomalies are also disallowed. According to G2, there can be effectively no cycle comprising anti-dependency edges between transactions. The expression

$$\neg (R^i(x)_{tx} FW^j(x)_{ty} FR^k(x)_{ty} \in St) \wedge (R^l(x)_{ty} FW^m(x)_{tx} FR^n(x)_{tx} \in St) \wedge (v'_k = v'_j) \wedge (v'_n = v'_m))$$

proscribes anti-dependency cycles. Thus, the ConSpec expression disallows G2, and effectively follows PL-3. Following the same line of reasoning, we can prove that the ConSpec expression in Equation 12 effectively specifies PL-2.99.

8. IMPLEMENTATION

We provide an automated verification tool built using the Spin model checker [9]. This tool can verify whether a given session trace satisfies a given consistency model or isolation level.

9. EQUIVALENCE BETWEEN DEPENDENCY RELATIONS AND LTL RELATIONS OF CONSPEC

State-of-the-art definitions specify consistency levels in terms of dependency relations, namely $w \xrightarrow{wr} r$, $w \xrightarrow{ww} w'$, and $w \xrightarrow{rw} r$ dependency, respectively [8]. The $w \xrightarrow{wr} r$ dependency can be translated to the form $\exists (W^i \in st) \wedge (R^j \in St)$

$(v^{i'} = v^{j'})$, where we represent the write and read operations w and r involved in the dependency relation with the propositional variable W^i and R^j , respectively. Denoting w , and w' with propositional variables W^i and W^j , the dependency relation $w \xrightarrow{ww} w'$ is translated into ConSpec form as follows.

$$\begin{aligned} & \exists (W^{i'} \in st) \wedge (W^{j'} \in St) \wedge (\exists (R^{l''} \in St) (v_l'' = v_j'')) \\ & \wedge ((v_l' = v_i') \vee ((W^{i''} FR^{l''} \in St) \wedge (\neg W^{i''} FW^{m''} FR^{l''} \in St)))) \end{aligned} \quad (13)$$

Similarly, denoting r , and w with propositional variables R^i and W^j , the $r \xrightarrow{rw} w$ dependency is translated into ConSpec form as follows.

$$\begin{aligned} & \exists (R^{i'} \in st) \wedge (W^{j'} \in St) \wedge (\exists W^{k''} \in st) \\ & (v_i' = v_k') \wedge (\exists R^{l''} \in St) (v_l' = v_j') \wedge ((v_k' = v_l') \vee \\ & ((W^{k''} FR^{l''} \in St) \wedge (\neg W^{k''} FW^{m''} FR^{l''} \in St))) \end{aligned} \quad (14)$$

Adya et al. define isolation levels in terms of $w \xrightarrow{wr} r$, $w \xrightarrow{ww} w'$, and $w \xrightarrow{rw} r$ dependencies among transactions [1]. The wr dependency relation of the dependency graph can be translated into ConSpec form

$$\exists (W_{tx}^i \in st) \wedge (R_{ty}^j \in St) (v^{i'} = v^{j'}). \text{ The } ww \text{ dependency}$$

relation between transactions tx and ty can be translated to

$$\begin{aligned} & \exists (W_{tx}^{i'} \in st) \wedge (W_{ty}^{j'} \in St) \wedge (\exists (R_{ty}^{l''} \in St) \\ & (v_l'' = v_j'')) \wedge ((v_l' = v_i') \vee ((W_{tx}^{i''} FR_{ty}^{l''} \in St) \wedge \\ & ((\neg W_{tx}^{i''} FW_{tx}^{m''} FR_{ty}^{l''} \in St) \\ & \wedge (\neg W_{tx}^{i''} FW_{ty}^{m''} FR_{ty}^{l''} \in St)))) \end{aligned} \quad (15)$$

The rw dependency relation between transactions tx and ty can be translated into ConSpec form

$$\begin{aligned} & \exists (R_{tx}^{i'} \in st) \wedge (W_{ty}^{j'} \in St) \wedge (\exists W_{tx}^{k''} \in st) \\ & (v_i' = v_k') \wedge (\exists R_{ty}^{l''} \in St) (v_l' = v_j') \wedge ((v_k' = v_l') \\ & \vee ((W_{tx}^{k''} FR_{ty}^{l''} \in st'') \wedge ((\neg W_{tx}^{k''} FW_{tx}^{m''} FR_{ty}^{l''} \in St) \\ & \wedge (\neg W_{tx}^{k''} FW_{ty}^{m''} FR_{ty}^{l''} \in St)))) \end{aligned} \quad (16)$$

10. REFERENCES

- [1] A. Adya, B. Liskov, and P. E. O'Neil. Generalized isolation level definitions. In *ICDE*, pages 67–78, 2000.
- [2] P. Bailis, A. Ghodsi, J. M. Hellerstein, and I. Stoica. Bolt-on causal consistency. In *SIGMOD '13*.
- [3] E. A. Brewer. Towards robust distributed systems (Invited Talk). In *Proc. of the 19th ACM SIGACT-SIGOPS Symposium on Principles of Distributed Computing*, 2000.
- [4] S. Burckhardt. Principles of eventual consistency. *Found. Trends Program. Lang.*, 1(1-2):1–150, Oct. 2014.
- [5] G. Chockler, R. Friedman, and R. Vitenberg. *Consistency Conditions for a CORBA Caching Service*, pages 374–388. Springer Berlin Heidelberg, Berlin, Heidelberg, 2000.
- [6] R. Ehlers and V. Raman. Low-effort specification debugging and analysis. In *Proceedings 3rd Workshop on Synthesis, SYNT 2014, Vienna, Austria, July 23-24, 2014.*, pages 117–133, 2014.
- [7] S. Gilbert and N. Lynch. Brewer's conjecture and the feasibility of consistent, available, partition-tolerant web services. *SIGACT News*, 33(2):51–59, June 2002.
- [8] J. L. Hennessy and D. A. Patterson. *Computer Architecture, Fifth Edition: A Quantitative Approach*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 5th edition, 2011.
- [9] G. Holzmann. *Spin Model Checker, the: Primer and Reference Manual*. Addison-Wesley Professional, first edition, 2003.
- [10] L. Lamport. The temporal logic of actions. *ACM Trans. Program. Lang. Syst.*, 16(3):872–923, May 1994.
- [11] S. Owens, S. Sarkar, and P. Sewell. A better x86 memory model: X86-tso. In *Proceedings of the 22Nd International Conference on Theorem Proving in Higher Order Logics, TPHOLs '09*, pages 391–407, Berlin, Heidelberg, 2009. Springer-Verlag.
- [12] D. B. Terry, A. J. Demers, K. Petersen, M. Spreitzer, M. Theimer, and B. W. Welch. Session guarantees for weakly consistent replicated data. In *Proceedings of the*

11. VIOLATION EXAMPLES: VERIFIED AGAINST CONSPEC SPECIFICATIONS

Using example session traces, we illustrate situations where a given consistency model is violated. Let us consider a session trace st_1 : $w^1(x, 1), w^2(x, 2), r^1(x)2, r^2(x)1$ to demonstrate how violation of RYW consistency can occur. From the session trace st_1 , the following LTL expression $W^1FW^2FR^1FR^2$ follows directly from the temporal order of execution as per the session trace. In st_1 , the read operation $r^1(x)2$ follows the write operation $w^1(x, 1)$ can be expressed by the LTL expression W^1FR^1 ; which satisfies the condition given in LHS of the shorter specification E^s in Equation 2 if $W^i = W^1, R^j = R^1$. The session trace st_1 can be transformed into the following partial legal serializations (that also comprises all write operations from other clients: $Si_p^1 = W^1FR^2FW^2FR^1, Si_p^2 = W^2FR^1FW^1FR^2$, and $Si_p^3 = R^1FW^2FW^1FR^2$. Only for Si_p^1 , the condition $W^1FR^2 \in Si_p^1$ holds, which does not hold for Si_p^2 and Si_p^3 . thus the condition corresponding to the RHS of the implies operator \rightarrow in E^s , given in Equation 4.1, holds. Next, we consider the sequence in the session trace where the write operation $w^1(x, 1)$ is followed by the read operation $r^2(x)1$, expressed by the LTL expression W^1FR^2 . As discussed before a partial serialization Si_p^1 exists for the session trace such that $W^1FR^2 \in Si_p^1$, satisfying the first condition in the RHS of Equation 4.1. Next, we consider the sequence write operation W^2 followed by R^1 in the session trace, given as W^2FR^1 . This implies that the condition W^iFR^j in the LHS of the E^s condition in Equation 4.1 holds, where $W^i = W^2$ and $R^j = R^1$. Again a legal partial serialization Si_p^1 exists such that $W^2FR^1 \in Si_p^1$. Next, we consider the sequence write operation W^2 followed by R^2 in the session trace, given as W^2FR^2 . This implies that the condition W^iFR^j in the LHS of E^s in Equation 4.1 holds, where $W^i = W^2$ and $R^j = R^2$. Si_p^1 is the only legal partial serialization that satisfies the condition $W^iFR^j \in Si_p$ in the RHS of Equation 4.1 for all the above write-read operation sequences. However, Si_p does not satisfy $W^iFR^j \in Si_p$ for the write-read sequence W^2FR^2 . For anomaly condition C in Equation 4.1 to hold, they cannot exist in the global session trace a read operation r^k preceding the read operation r^j (which is $r^2(x)1$) and a write operation w^l preceding w^i (which is $w^2(x, 2)$ in this case) such that the LTL condition $(v^{k'} = v^{i'}) \wedge ((v^{j'} = v^{l'}) \vee ((W^{l'}FR^{j'} \in St) \wedge (\neg W^{l'}FW^{m'}FR^{j'} \in St)))$ hold. However, in the given session trace st_1 , there does exist a read $r^1(x)2$ and a write $w^1(x, 1)$, such that conditions $(v^{k'} = v^{i'})$ and $(v^{j'} = v^{l'})$ hold, thus, in turn, violating the anomaly condition. Hence, the RYW condition is violated for the given session trace.

Let us consider a session trace st_2 : $w^1(x, 1), st_2': w^2(x, 2)$, and $st_2'': r^1(x)1, r^2(x)2, r^3(x)1$ to illustrate how violation of MR consistency can occur. Since the above session traces comprise write operations followed by read on the same object, as well as multiple read operations on same object x , they satisfy the conditions given in LHS of Equation 2. First, let us consider the sequence of read operations $r^1(x)1$ and $r^2(x)2$ in st_2'' , given by the LTL expression R^1FR^2 . If and only if $R^i = R^1$ and $R^j = R^2$, the above sequence of reads can satisfy the LHS of the condition E^s in Equation 2. The above traces can be transformed into legal partial serialization order $Si_p = W^1FR^1FR^3FW^2FR^2$, such that

$R^1FR^2 \in Si_p$ holds, thus satisfying the first condition the RHS of the implies operator \rightarrow in E^s in Equation 2. Next, we consider the sequence $r^1(x)1$ followed by $r^3(x)1$ in the above traces, given as the LTL expression R^1FR^3 . The above condition satisfies the LHS of E^s Equation 2, given that $R^i = R^1$ and $R^j = R^3$, which, in turn, implies $v_i = v_1 = 1$ and $v_j = v_2 = 1$. As discussed above, a legal partial serialization Si_p exists for the above trace, such that $R^1FR^3 \in Si_p$ holds; the condition E^s in the RHS of Equation 2 holds for the above sequence of reads. Next, we consider the sequence $r^2(x)2$ followed by $r^3(x)1$ in the above traces, given as R^2FR^3 . However, the above sequence of reads does not match the order among these reads in the serialization Si_p , i.e., $R^2FR^3 \notin Si_p$ holds; thus it does not satisfy the first condition in the RHS of Equation 2. Further, for the LHS of Equation 2 to hold for the above sequence of reads, $R^i = R^2$ and $R^j = R^3$, which, in turn, implies $v_i = v_2 = 2$ and $v_j = v_3 = 1$. For the anomaly condition C in Equation 2 to hold for the sequence of read operations $r^1(x)1$ followed by $r^3(x)1$, there cannot exist write operations w^m and w^n corresponding to the reads r^i and r^j such that the LTL condition $(v^{i'} = v^{n'}) \wedge ((v^{j'} = v^{m'}) \vee ((W^{m'}FR^{j'} \in St) \wedge (\neg W^{m'}FW^{p'}FR^{j'} \in St)))$ holds. Here, $W^m = W^1$ and $W^n = W^2$, since W^m and W^n are the only write operations in the given session history. Hence, $v_m = v_1 = 1$ and $v_n = v_2 = 2$. Also, it is given that $v_i = v_1 = 1$ and $v_j = v_2 = 2$. Hence, the conditions $(v^{i'} = v^{n'})$ and $(v^{j'} = v^{m'})$ is satisfied; thus, in turn, the above LTL condition is also satisfied. Thus, the violation condition C in Equation 2 holds. Hence, the above traces violate the MR condition.

Let us consider a session trace st_3 : $w^1(x, 1), r^1(x)2, st_3': r^2(x)1w^2(x, 2), r^3(x)1$, and $st_3'': r^4(x)2, r^5(x)1$ to demonstrate a violation of Causal consistency. Since, there are successive operations with direct precedence relation on the object x , st_3 satisfies LHS of the E^s condition in Equation 3. We can transform the above session traces into a partial serialization $Si_p = W^1FR^2FR^3FR^5FW^2FR^1FR^4$. First, let us consider the sequence of operations $w^1(x, 1)$ and $r^1(x)2$ in the above traces, given by the LTL expression W^1FR^1 . The above expression satisfies the condition $Op^iFR^jOp^j$, where $Op^i = W^1$ and $Op^j = R^1$; thus satisfying the LHS of the implies operator \rightarrow in E^s in Equation 3. The above sequence of writes and reads satisfies the condition $W^1FR^1 \in Si_p$, thus satisfying the first condition the RHS of the implies operator \rightarrow in E^s in Equation 3. Following the same line of argument, the sequence of operations $w^1(x, 1)$ and $r^2(x)1$, $w^1(x, 1)$ and $r^3(x)1$, $w^1(x, 1)$ and $r^4(x)2$, $w^1(x, 1)$ and $r^5(x)1$ and $w^2(x, 2)$ satisfy the LHS and RHS of E^s in the Equation 3. Next, let us consider the sequence of operations $w^2(x, 2)$ and $r^3(x)1$, which is expressed using LTL condition W^2FR^3 . Thus, the above sequence satisfies LHS of E^s in the Equation 3. The above sequence of writes and reads do not satisfies the condition $W^2FR^3 \in Si_p$, violating the RHS of E^s in Equation 4.1. Let us consider the anomaly condition C in Equation 3 for the write and read operations $w^2(x, 2)$ (which is w^i in this case) and $r^3(x)1$ (which is r^j in this case). Since $w^2(x, 2)$ is followed by $r^3(x)1$ in the session trace, they satisfy the LHS of the condition C . For the RHS of C to hold, there cannot exist write operations w^m preceding the write $w^2(x, 2)$ such that the LTL condition $(v^{j'} = v^{m'}) \vee ((W^{m'}FOP^{j'} \in st) \wedge (\neg W^{m'}FW^{n'}FOP^{j'} \in st))$ holds. We can assign $W^m = W^1$ since $w^1(x, 1)$ precedes $w^2(x, 2)$. Subsequently the condition $(v^{j'} = v^{m'})$ holds (since $R^j = R^3$ and $v_j = 1$ in this case). Hence, the RHS of C is satisfied for Equation 3 for the above sequence of operations.

It results in violation of Causal consistency condition.

All the above session traces violate the specification of Strict Serializability, given in Equation 4. Consider the session traces st_3 , st'_3 , and st''_3 , which can be transformed into the legal serialization Si , which is, in this case, equivalent to Si_p in absence of writes from other client applications in the given session history (see above). The sequence of operations $w^1(x, 1)$ and $r^1(x)2$ in the above session trace satisfies the LHS of the condition E^s (i.e., the condition $Op^i X Op^j$). Since $v_i \neq v_j$, the above sequence of write and read operations do not satisfy the condition in the RHS of E^s , i.e., it violates the condition $W^1 X R^1 \in Si$. For the anomaly condition C to hold for the above sequence of operations, let us consider the sequence of operations $w^2(x, 2)$ and $r^3(x)1$. For the RHS of C to hold, there cannot exist write operations r^k following the read $r^3(x)1$ such that the LTL condition $((v^{k''} = v^{i'}) \wedge (v^{j'} \neq v^{i'})) \wedge \neg \exists ((Op^{i'} = W^{i'} \oplus R^{i'}) \wedge (Op^{i'} FW^{j'} FR^{k'} \in st) \wedge (Op^{i'} FW^{j''} FR^{k'} FR^{l'} \in St))$ holds. We can assign $r^k = r^4$ since $r^4(x)2$ follows $r^3(x)1$. Subsequently the conditions $(v^{k''} = v^{i'})$ and $(v^{j'} \neq v^{i'})$ holds, since r^j yields $v_j = 1$, and r^k yields $v_k = 2$. Thus, it also violates the anomaly condition in Equation 4, violating strict serializability.

We consider the session traces $st_4: w^1(x, 1)$, $st'_4: r^1(x)1, w^2(x, 2)$, and $st''_4: r^2(x)2, r^3(x)1$ to show how WFR consistency can be violated. The above session traces can be transformed into a legal serialization $Si_p = W^1 FR^1 FR^2 FW^{2i} W^{2j}$. The write operation $w^1(x, 1)$ follows the read operation $r^1(x)1$, and can be expressed as $W^1 F^x R^1$. The above expression satisfies the condition $W^i FR^j$, where $W^i = W^1$ and $R^j = R^1$. Hence, the above traces satisfy the condition in the LHS of Equation 5. Further, by observation, the above sequence of write-read operations satisfies the condition $W^1 FR^1 \in Si_p$. Thus, the condition E^s in the RHS of Equation 5 holds. According to the anomaly condition C in Equation 5, there must not exist a read operation $r^n(x)$ such that the conditions $(R^i F^x W^j F^x R^n \in st) \wedge (W^j F^x R^n \in st)$ and $v^n = v^i$ are satisfied. Since $r^2(x)1$ succeeds $w^2(x, 2)$, we can assign $R^n = R^2$. But, the read operation $r^2(x)1$ returns a value 1 which was written by $w^1(x, 1)$; thus the condition $v^n = v^i$ holds. Thus, st_4 violates the second condition in RHS of Equation 5; hence, it violates WFR consistency.

Let us consider the session traces $st_5: w^1(x, 1), r^1(x)2, w^2(x, 2)$, $st'_5: r^2(x)2$, and $st''_5: r^3(x)1, r^4(x)2$, and demonstrate violation of MW consistency with it. The above session traces can be transformed in a legal serialization $Si = W^1 FW^2 FR^1 FR^2$. Since the above session traces comprise a sequence of successive writes $w^1(x, 1), w^2(x, 2)$ on same object x , it can be expressed as $W^i FW^j$, thus satisfying the LHS of E^s in Equation 6, where $W^i = W^1$ and $W^j = W^2$. Also, the above sequence of write operations satisfies the first condition of the RHS of Equation 6, i.e., $W^i FW^j \in Si$. The above traces can be transformed into a legal serialization order $Si = w^1(x, 1), r^2(x)1, w^2(x, 2), r^1(x)2$, thus satisfying the first condition in the RHS of Equation 5. For the anomaly condition C to hold, there cannot exist a sequence of write operations $w^1(x, 1)$ and $w^2(x, 2)$ followed by a sequence of read operations r^k and r^l such that the following condition holds: $(\neg R^{k'} FR^{l'} \in St) \wedge (v_k' = v_j') \wedge (v_l' = v_i')$. For the above session trace, we assign the operations $r^2(x)2$ and $r^3(x)1$ to the read operation r^k and r^l . However, the

above read operations r^k yields $v_k = 2$, and r^l yields $v_l = 1$, such that the condition $(v_k' = v_j') \wedge (v_l' = v_i')$ holds.

Consider a group of 2 clients (or processors), namely Cl_1 and Cl^2 , respectively. Let us consider that operations are invoked from the above 2 clients, and a session trace st_6 is collected as follows. $st_6: w^1(x, 1), r^1(y)1, st'_6: w^2(y, 1), r^2(x)2$. For the above session trace to satisfy TSO consistency, the condition $(W^{i'}(x) FR^{j''}(x) \in St)$

$(v_j'' \neq v_i') \wedge (R^{k''}(x) FR^{j''}(x) \in St) \wedge (v_k'' \neq v_i') \wedge (R^{k'} FR^{j'} \in st) \wedge (W^{l'} FW^{i'} \in St) \wedge (v_k' = v^{i'}) \wedge ((v^{j'} = v^{l'}) \vee ((W^{l'} FR^{j'} \in St) \wedge \neg W^{m'} (W^{l'} FW^{m'} FR^{j'} \in St)))$ must hold. In the above session traces, let us assign the operation $w^1(x, 1)$ as w^i . $r^2(x)2$ occurs in another session trace, and thus invoked was from a different client. Further, $r^2(x)2$ follows $w^1(x, 1)$ in the global session order. Hence we can assign the operation $r^2(x)2$ as r^j . Since $r^2(x)2$ returns 2 instead of 1, the condition $v_j'' \neq v_i'$ does not hold. Also, there doesn't exist an operation r^k preceding $r^2(x)2$ that satisfies the condition $v_k'' \neq v_i'$. Thus, the given session trace violates Equation 7.

Consider two client applications (or processors), namely Cl_1 and Cl^2 , respectively. Consider a session trace collected by executing operations on the above client applications as follows. $st_7: w^{2i}(x, 1), w^2(y, 1)$, $st'_7: r^1(x)1, st''_7: r^2(y)1, r(x)1$. To satisfy the PC consistency, any sequence of write operations w^i and w^j invoked by a client must satisfy the condition $(\neg (W^{i'} FW^{j'} \in st) \wedge (st' \in St) \wedge (R^{k'} FR^{l'} \in st')) ((v_k'' = v_j') \wedge (v_l' = v_i'))$. Thus, there must not exist a sequence of reads r^k and r^l such that the condition $(v_k'' = v_j') \wedge (v_l' = v_i')$ holds. In the above session traces, the write operations are executed according to the precedence order $w^1(x, 1)$ followed by $w^2(y, 1)$, hence we assign $w^1(x, 1)$ and $w^2(y, 1)$ as w^i and w^j . Also, let us assign the sequence of read operations $r(x)1$ and $r^2(y)1$ as r^k and r^l , respectively. According to the above condition, reads against the above writes should view the results of the above write operations according to the invocation order of these writes. However, in the given session trace, the client application Cl_2 observes the above writes out of order. Since $r^2(y)1$ executes before $r(x)1$ in the above session traces, Cl_2 observes the result of $w^2(y, 1)$ before observing result of $w^1(x, 1)$. Thus, the above traces violate Equation 8.

Consider a session trace $st_8: w_{tx}^1(x, 2), w_{ty}^2(x, 5), w_{ty}^1(y, 5), c_{ty}, w_{tx}^2(y, 8), r_{ty}^1(x)5, r_{ty}^2(y)8, c_{tx}$. According to the specification of PL-1 isolation level (refer to Equation 9), there cannot exist read operations r_{ty}^m and r_{ty}^n such that they satisfy the condition $((v_n' = v_i') \wedge (v_m' = v_i')) \vee ((v_n' = v_j') \wedge (v_m' = v_k'))$. In the session trace st_8 , the write operations $w_{tx}^1(x, 2), w_{tx}^2(y, 8), w_{ty}^2(x, 5)$, and $w_{ty}^1(y, 5)$ are $W_{tx}^i, W_{tx}^j, W_{ty}^k$, and W_{ty}^l , respectively. Also, we can assign the read operations $r_{ty}^1(x)5$ and $r_{ty}^2(y)8$ as r_{ty}^m and r_{ty}^n , respectively. Under the above assignment, we have $v_m = 8$ and $v_n = 5$, which in turn, satisfies the conditions $(v_n' = v_j') \wedge (v_m' = v_k')$. This, in turn, satisfies the anomaly condition for PL-1. Thus, st_8 violates PL-1.

Let us consider a session trace $st_9: w_{tx}^1(x, 1), r_{ty}^1(x)1, w_{tx}^2(x, 2), c_{ty}, c_{tx}$. According to PL-2

(refer to Equation 10), a transaction never must read a value from an operation if the transaction executing above operation has not yet committed. In other words there must not exist a read operation $r^k(x)_{ty}$ such that $r^k(x)_{ty}$ reads a value written by a write operation $w^i(x)_{tx}$ that does not write the final (committed) version of the object x for the transaction tx , i.e., the condition $\neg((W^i(x)_{tx}'FW^j(x)_{tx}'Fc_{tx} \in st) \wedge (v'_k = v'_i)) \wedge \neg((W^i(x)_{ty}'FW^j(x)_{ty}'Fc_{ty} \in st) \wedge (v'_k = v'_i))$

does not hold. In the session trace st_9 , the read operation $r^1_{ty}(x)1$ returns a value 1 written by the operation $w^1_{tx}(x, 1)$ in transaction tx . However, the write operation $w^2_{tx}(x, 2)$, which follows $w^1_{tx}(x, 1)$ in st_9 , writes the committed version (i.e., the final version) of x , which is 2. Thus, the above condition does not hold for st_9 , resulting in violation of PL-2.

Consider a session trace

st_{10} : $r^1_{tx}(tx)50, r^1_{ty}(x)-50, r^2_{ty}(y)100, w^1_{ty}(x, 100), w^2_{ty}(y, 50), c_{ty}, w^1_{tx}(y, -50), r^2_{ty}(x)100, r^2_{tx}(y)-50, c_{tx}$. In st_{10} , the read operation $r^1_{tx}(x)50$ on object x in tx is followed by write operation $w^1_{ty}(x, 100)$. The read operation $r^2_{ty}(x)100$ following the commit c_{ty} returns the value 100, indicating a rw dependency between transactions tx and ty . Similarly, the read operation $r^2_{ty}(y)100$ on object y in ty is followed by write operation $w^1_{tx}(y, 50)$. The read operation $r^2_{tx}(y)-50$ following the commit c_{ty} returns the value -50, implying a read-write dependency between ty and tx . Thus, there is a read-write dependency cycle in S_{11} , violating PL-3 specifications. Let us consider a session trace st_{11} : $r^1_{tx}(tx)50, r^1_{ty}(x)-50, r^2_{ty}(x)100, w^1_{ty}(x, 100), w^2_{ty}(y, 50), c_{ty}, w^1_{tx}(x, -50), r^2_{ty}(x)100, r^2_{tx}(x)-50, c_{tx}$. Following similar line of reasoning, we can observe that the session trace st_{11} violates the isolation level PL-2.99, given in Equation 12.